

# Lung Cancer Detection Milestone

Jay DeStories, Jason Fan, Alex Tong

March 2017

## 1 Introduction

In this document, we record the intermediary results for the Multi-Instance Network Lung Cancer Detection project.

Please refer to project proposal for references to the proposed pipeline and techniques.

### 1.1 Things we didn't expect to be so difficult

Our team found the following two problems more difficult than expected:

1. Getting working implementation of AlexNet [1] and learning TensorFlow
2. Pre-processing Data and working with DICOM images

### 1.2 Problems with Preprocessing

With the LUNA and Kaggle Dicom data amounts to almost a terabyte of storage, we made some mistakes with data processing and had to process the Kaggle Data twice.

We had trouble converting Hounsfield Units (hu), the unit of measurement for DICOM slice densities, to a sensible numeric value a traditional image classification network would be able to consume.

Hounsfield Units measure density and map different parts of the human bodies to a high range of values, we experimented with the DICOM slices to find appropriate thresholded minimum and maximum values so that resulting images retained enough contrast. You can see in Figure 2, that the lung scaled to reasonable pixel values has very little contrast. In figure 4 you can see our rescaling and thresholding efforts to improve the contrast of the image. We reason that the closer we can get the histogram of the lung to look to a typical image in the ImageNet dataset, then the better VGG will perform in discerning features on the lung image.

We attempted to scale the image such that the lung had a black background and the majority of the lung pixels were in the center range of pixel intensities. We speculate that it might be advantageous to increase the contrast further (essentially stretch the bump in

the middle of Figure 3) so that our image histogram might look more like an ImageNet histogram.




Figure 1: Histogram of Pixel Values of a resized interpolated lung




Figure 2: Lung Slice without preprocessing




Figure 3: Histogram of Pixel Values of a scaled resized interpolated lung




Figure 4: Histogram of Pixel Values of a scaled thresholded resized interpolated lung

### 1.3 From AlexNet in Keras to VGG in TensorFlow

The Multi-Instance Network from the University of California Irvine (UCI) researchers was implemented with a custom, unmaintained version of Keras that we found, albeit too late, to be untenable.

We decided to move to using TensorFlow for clearer documentation and more support from the online community. At first, we found TensorFlow difficult to understand, but we eventually found a simple enough implementation of VGG to work with. (From: [github.com/machrisaa/tensorflow-vgg](https://github.com/machrisaa/tensorflow-vgg))

## 2 Pipeline Progress

We have implemented the following:

1. Memory concious RGB CNN to Grayscale CNN conversion
2. Feature extraction with pre-trained VGG-16
3. Feature extraction with pre-trained VGG-19
4. Multi-Instance 2-class logistic regression
5. Boosted decision tree

### 2.1 RGB to Grayscale conversion

The Keras implementation of Multi-Instance Network feeds the first convolutional an image of shape (224,224,3) in which the black and white dicom image is copied once onto each Blue, Red and Green channel.

We attempted to mimic this technique with the stock VGG implementation to extract slice-wise image features from the last fully convolutional layer and found that we could not fit an entire lung volume into VRAM.

We noticed that for weight vector  $w$  along the channels axis of a convolution, the result of the convolution with pixel vector  $x$  where each channel has value  $\hat{x}$ , channelwise mean  $\mu$  and bias term  $b$  was,

$$w*(x-\mu)+b = (w*x)+(b-w*\mu) = \hat{x}\|w\|_1+(b-w*\mu)$$

Which meant that we could eliminate the redundant 3-channel input dimension and make any convolutional neural network effectively consume grayscale images by manipulating pretrained weights and biases for the first layer. We simply sum the convolution along the axis of the input channels dimension and subtract the term  $(w * \mu)$  from the bias.

We did this for both VGG-16 and VGG-19 and were able to forward pass entire lung volumes to extract slice-wise features.

### 2.2 VGG feature extraction

We perform the above mentioned RGB to Grayscale conversion for both VGG-16 and VGG-19 and were able to reduce VRAM usage and forward pass entire lung volumes to extract slice-wise features [2]. A feature tensor of shape (7, 7, 512) is extracted for each slice, where each of the 49 receptive fields is represented by a feature vector of length 512.

The resulting bag of features that is then consumed by the Multi Instance 2-class Logistic Regression and Boosted Decision Tree Models, is a tensor of shape (60, 7, 7, 512).

### 2.3 Multi Instance 2-class logistic regression

Let  $F$  be a set of  $N$  feature vectors that represent the instances of a given lung slice. Then  $r$ , the vector of activations, is defined element wise by

$$r_i = w^T f_i + b$$

. Where  $f_i$  and  $r_i$  is the  $i$ -th element of  $\mathbf{r}$  and  $F$  respectively.

We then use  $\mathbf{r}$  to predict the probability of cancer  $p(y = 1)$  where

$$t = p(y = 1) = \max_{r_i} \sigma(r_i)$$

The authors of the Multi-Instance learning paper proposed the following per example loss functions to regularize hyper-parameters at training time [3].

Normal softmax loss:

$$L_{max} = -\log(t^y + (1 - t)^{(1-y)})$$

Sparse softmax loss:

$$L_{max} = -\log(\lambda_t(t^y + (1 - t)^{(1-y)}) + \lambda_r \|\mathbf{r}\|)$$

## 2.4 Intermediate Results with Boosted decision tree with XGBoost

As a preliminary evaluation of the features extracted using VGG, we trained gradient boosted decision trees on the VGG-16 and VGG-19 features. One challenge we discovered is that the features take up too much space to fit the entire dataset into memory. For the following results, we loaded the first 500 lungs for training and used 100 lungs for validation. We used a max depth of 10 and recorded the final error, area under curve, and log loss on the validation set for each set of input features.

Features	Error	AUC	Log Loss
VGG-16	0.25	0.521739	0.611064
VGG-19	0.26	0.440429	0.672109

In all cases the boosted decision trees reached perfect error and area under curve within a few rounds. It's clear from our results that this model overfits significantly, but this provides a good baseline for further work.

## References

- [1] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E Hinton. Imagenet classification with deep convolutional neural networks. In F. Pereira, C. J. C. Burges, L. Bottou, and K. Q. Weinberger, editors, *Advances in Neural Information Processing Systems 25*, pages 1097–1105. Curran Associates, Inc., 2012.
- [2] Karen Simonyan and Andrew Zisserman. Very deep convolutional networks for large-scale image recognition. *arXiv preprint arXiv:1409.1556*, 2014.

- [3] Wentao Zhu, Qi Lou, Yeeleng Scott Vang, and Xiaohui Xie. Deep multi-instance networks with sparse label assignment for whole mammogram classification. *CoRR*, abs/1612.05968, 2016.