

Lung Cancer Detection Milestone

Jay DeStories, Jason Fan, Alex Tong

March 2017

1 Introduction

In this document, we record the intermediary results for the Multi-Instance Network Lung Cancer Detection project.

Please refer to project proposal for references to the proposed pipeline and techniques.

1.1 Things we didn't expect to be so difficult

Our team found the following two problems more difficult than expected:

1. Getting working implementation of AlexNet and learning TensorFlow
2. Pre-processing Data and working with DICOM images

1.2 Problems with Pre-processing

With the LUNA and Kaggle Dicom data amounts to almost a terabyte of storage, we made some mistakes with data processing and had to process the Kaggle Data Twice.

We had trouble converting Hounsfield Units (hu), the unit of measurement for DICOM slice densities, to a sensible numeric value a traditional image classification network would be able to consume.

Hounsfield Units measure density and map different parts of the human bodies to a high range of values, we experimented with the DICOM slices to find appropriate thresholded minimum and maximum values so that resulting images retained enough contrast.

INSERT HISTOGRAM OF HU UNITS HERE

INSERT DICOM LUNG without morph closing with low contrast...

INSERT Pre-Processed image with segmented lung and thresholded values

1.3 From AlexNet in Keras to VGG in TensorFlow

The Multi-Instance Network from the University of California Irvine (UCI) researchers was implemented with a custom, unmaintained version of Keras that we found, albeit too late, to be untenable.

We decided to move to using TensorFlow for clearer documentation and more support from the online community. At first, we found TensorFlow difficult to understand, but we eventually found a simple enough implementation of VGG to work with. (From: github.com/machrisaa/tensorflow-vgg)

2 Pipeline Progress

We have implemented the following:

1. Memory conscious RGB CNN to Grayscale CNN conversion
2. Feature extraction with pre-trained VGG-16
3. Feature extraction with pre-trained VGG-19
4. Multi-Instance 2-class logistic regression
5. Boosted decision tree

2.1 RGB to Grayscale conversion

The Keras implementation of Multi-Instance Network feeds the first convolutional an image of shape (224,224,3) in which the black and white dicom image is copied once onto each Blue, Red and Green channel.

We attempted to mimic this technique with the stock VGG implementation to extract slice-wise image features from the last fully convolutional layer and found that we could not fit an entire lung volume into VRAM.

We noticed that for weight vector w along the channels axis of a convolution, the result of the convolution with pixel vector x where each channel has value \hat{x} , channelwise mean μ and bias term b was,

$$w*(x-\mu)+b = (w*x)+(b-w*\mu) = \hat{x}\|w\|_1+(b-w*\mu)$$

Which meant that we could eliminate the redundant 3-channel input dimension and make any convolutional neural network effectively consume grayscale images by manipulating pretrained weights and biases for the first

layer. We simply sum the convolution along the axis of the input channels dimension and subtract the term $(\mathbf{w} * \boldsymbol{\mu})$ from the bias.

We did this for both VGG-19 and VGG-19 and were able to forward pass entire lung volumes to extract slice-wise features.

2.2 VGG feature extraction

ADD VGG AS A CITATION

We do the abovementioned RGB to Grayscale conversion for both VGG-19 and VGG-19 and were able to reduce VRAM usage and forward pass entire lung volumes to extract slice-wise features. A feature tensor of shape (7, 7, 512) is extracted for each slice, where each of the 49 receptive fields is represented by a feature vector of length 512.

The resulting bag of features that is then consumed by the Multi Instance 2-class Logistic Regression and Boosted Decision Tree Models, is a tensor of shape (60, 7, 7, 512).

2.3 Multi Instance 2-class logistic regression

ADD MIL CITATION

Let F be a set of N feature vectors that represent the instances of a given lung slice. Then \mathbf{r} , the vector of activations, is defined by elementwise by

$$r_i = \mathbf{w}^T \mathbf{f}_i + b$$

. Where f_i and r_i is the i -th element of \mathbf{r} and F respectively.

We then use \mathbf{r} to predict the probability of cancer $p(y = 1)$ where

$$t = p(y = 1) = \max_{r_i} \sigma(r_i)$$

The authors of the Multi-Instance proposed the following per example loss functions to regularize hyperparameters at training time.

Normal softmax loss:

$$L_{max} = -\log(t^y + (1 - t)^{(1-y)})$$

Sparse softmax loss:

$$L_{max} = -\log(\lambda_t(t^y + (1 - t)^{(1-y)}) + \lambda_r \|\mathbf{r}\|)$$

2.4 Boosted decision tree with XG-Boost

3 Intermediary results